Programming Massively Parallel Processors

CH1: Introduction

• multicore CPUs: maximize execution speed of sequential programs
• many-core GPUs: execution throughput of parallel applications
• large performance gap between parallel and sequential execution

Why?

• developers move computationally intensive parts of software to GPUs

• design of a CPU is optimized for sequential code performance

• memory bandwidth is another issue. GPUs move data much faster in and out of its DRAM.

• video games require massive number of floating point calculations per video frame, being executed in parallel and GPUs have been used for this purpose.

CUDA (Compute Unified Device Architecture)

• programming model created by NVIDIA to support joint CPU/GPU execution of an application.

• CUDA-capable GPU is organized into an array of highly threaded streaming multiprocessors (SMs). SMs combine to form a building block. SMs have a number of streaming processors (SPs) that share control logic and instruction cache.
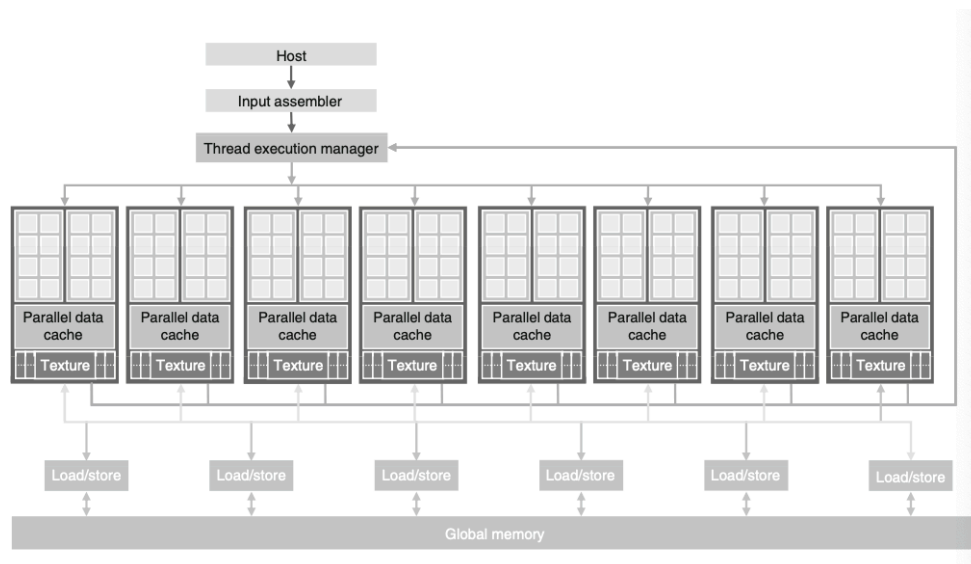
• To experience speedup offered by parallelization, a large part of the application's execution time must be in the parallel portion.

- Certain applications have portions better suited to CPUs and hence a combined CPU/GPU parallel computing capability is required. This is precisely what CUDA promotes.
- Key steps in parallel computing:
  ‣ identifying parts of application programs to be parallelized
  ‣ isolating the data to be used by the parallelizing code by using API functions to allocate memory on the parallel computing device
  ‣ using API functions to transfer data to the parallel computing device
  ‣ developing kernel functions that will be executed by individual threads in the parallel part
  ‣ launching kernel functions for execution by parallel parts
  ‣ transferring data back to the host processor with API function calls

CH2: History of GPU Computing

GPGPU
General Purpose Computing on GPUs.
- GPU processor array and frame buffer memory were designed to process graphics data and were too restrictive for general numerical applications.
- writes were extremely difficult → could only be emitted as a pixel color value and configure the frame buffer operation to write.
- the handful of useful applications created with general computations on a GPU → this field was called GPGPU.

GPU Computing
- NVIDIA developed Tesla GPU Architecture.
- programming paradigm to think of GPU like a processor.
- programming approach involved explicit declaration of data-parallel aspects of their workload.
- no longer need to use graphics API to access parallel computing capabilities.

CH3: Introduction to CUDA

Data Parallelism

- computing system consists of host (CPU) & devices (massively parallel processors)
- CUDA devices accelerate execution of applications by harvesting a large amount of data parallelism.
- Matrix multiplication $P = M \times N$:
  - ‣ as every entry p_{ij} is independent of each other, a large amount of data parallelism can be performed.

CUDA Program Structure

- CUDA program comprises phases that are executed either by the host (CPU) or a device such as a GPU.
- CUDA program is a unified source code comprising both host & device code.
  - ‣ NVIDIA C compiler (nvcc): host code = ANSI C; device code = ANSI C extended with keywords for data-parallel functions, called kernels.
  - ‣ kernel functions generate a large number of threads to exploit data parallelism.
  - ‣ when no device is available, one can execute the kernel on a CPU using the CUDA SDK or MCUDA tool.
- CUDA threads are faster to generate and schedule than CPU threads due to efficient hardware support.
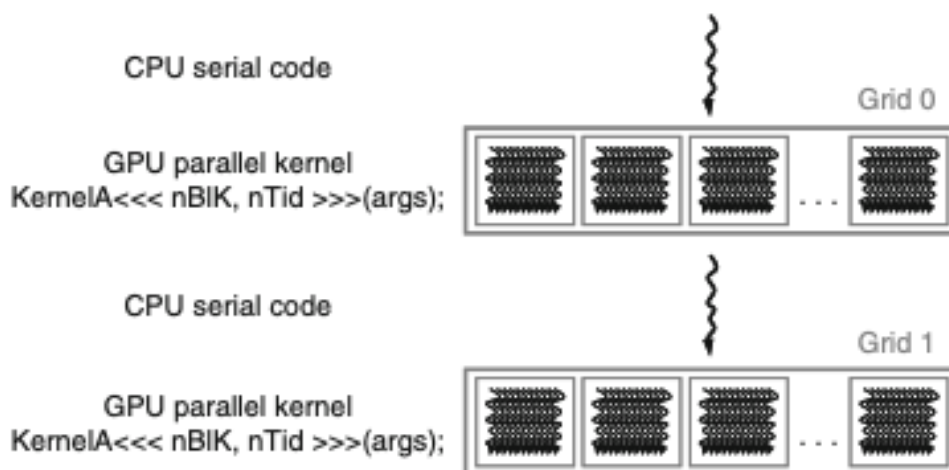
- CUDA program:

- starts with CPU execution; when a kernel function is invoked, execution is moved to the device.
- large numbers of threads are generated to take advantage of data parallelism, collectively called grids.
- when all threads finish execution, the corresponding grid is terminated and execution moves back to the host.

MATMUL Example

```
int main (void) {
    // 1. Allocate and initialize matrices M, N, P
    //    I/O to read input matrices M and N

    // 2. M * N on the device
    //    MatrixMultiplication(M, N, P, Width);

    // 3. I/O to write output matrix P
    //    Free matrices M, N, P

    return 0;
}
```
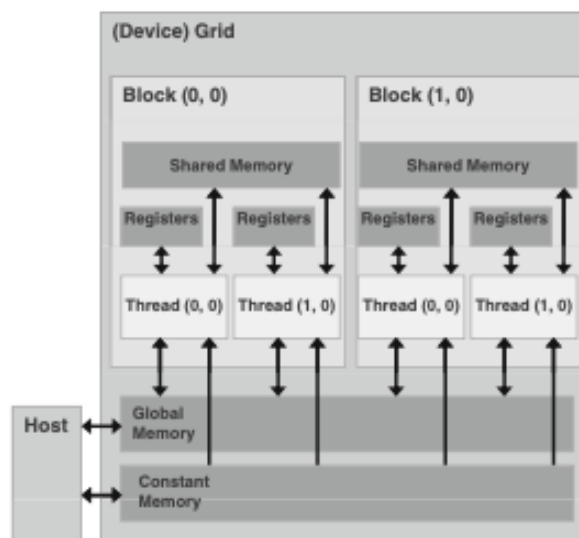
Device Memory & Data Transfer

CUDA runtime system provides API functions to perform memory allocation and data transfer between host and devices.
CUDA devices comprise global memory and constant memory; these are accessible from host code.
Constant memory is read-only for device code.

API functions cudaMalloc() and cudaFree() allocate and free global memory.
API function cudaMemcpy() transfers data between host & device memory.

- cudaMalloc()
  - Allocates object in the device global memory
  - Two parameters
    - **Address of a pointer** to the allocated object
    - **Size of** of allocated object in terms of bytes
- cudaFree()
  - Frees object from device global memory
    - Pointer to freed object
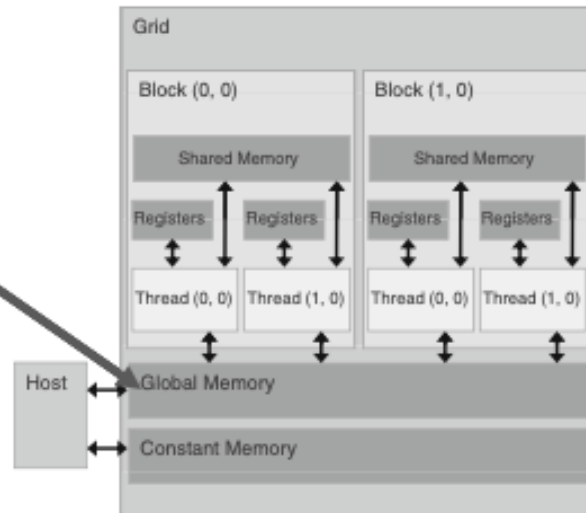
For cudaMalloc():
First parameter is a generic pointer (void)
Second parameter is the size in bytes
Example:

```
float *Md;

int size = Width * Width * sizeof(float);

cudaMalloc((void**)&Md, size);


// ...

cudaFree(Md);
```
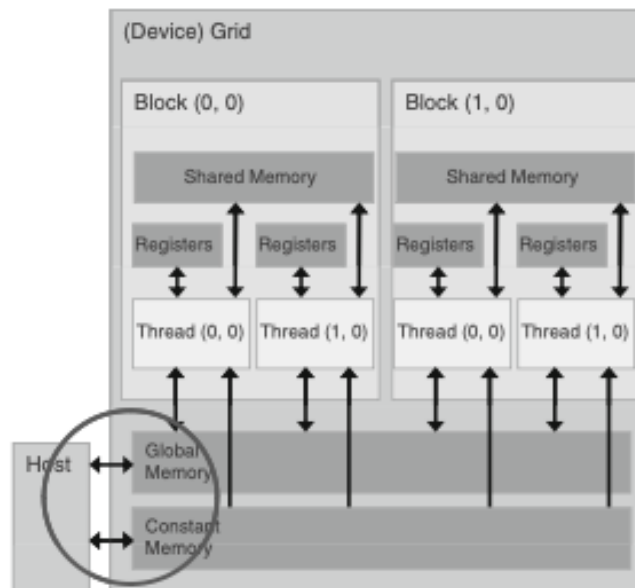
For cudaMemcpy():
First argument = destination pointer
Second = source pointer
Third = number of bytes
Fourth = direction (host→device, device→host, device→device)
Note: cudaMemcpy() cannot be used for memory transfer in multi-GPU systems.

- cudaMemcpy()
  - **Memory** data transfer
  - Requires four parameters
    - Pointer to destination
    - Pointer to source
    - Number of bytes copied
    - Type of transfer
      - Host to Host
      - Host to Device
      - Device to Host
      - Device to Device
  - Transfer is asynchronous

Example:
```
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
```
In the MatMul example, the main() function calls MatrixMultiplication(). MatrixMultiplication() allocates device memory, performs data transfers, and invokes the kernel that computes the result. This type of host-side function is called a stub function.

```
void MatrixMultiplication(float *M, float *N, float *P, int Width) {
    int size = Width * Width * sizeof(float);
    float *Md, Nd, Pd;

    // Allocate device memory for M, N, P
    cudaMalloc((void**)&Md, size);
    cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Nd, size);
    cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);
    cudaMalloc((void**)&Pd, size);

    // Kernel invocation (not shown)
    // ...

    // Copy P & free device memory
    cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
    cudaFree(Md); cudaFree(Nd); cudaFree(Pd);
}
```

Kernel Functions and Threading

In CUDA, a kernel function is executed by many threads in parallel. CUDA programming follows the single program, multiple data (SPMD) model.

	Executed on the:	Only callable from the:
__device__ float DeviceFunc()	device	device
__global__ void KernelFunc()	device	host
__host__ float HostFunc()	host	host

Device functions cannot include recursion or indirect calls.
A function can be annotated for both host & device, generating two versions.
Variables like threadIdx.x and threadIdx.y give thread coordinates.
Thread notation: $\mathrm{thread}_{\{\mathrm{threadIdx}.x,\ \mathrm{threadIdx}.y\}}$
Threads are organized hierarchically:
A grid contains one or more blocks
Each block has a unique 2D coordinate

- A thread block is a batch of threads that can cooperate with each other by:
  — Synchronizing their execution
    - For hazard-free shared memory accesses
  — Efficiently sharing data through a low-latency shared memory
- Two threads from two different blocks cannot cooperate
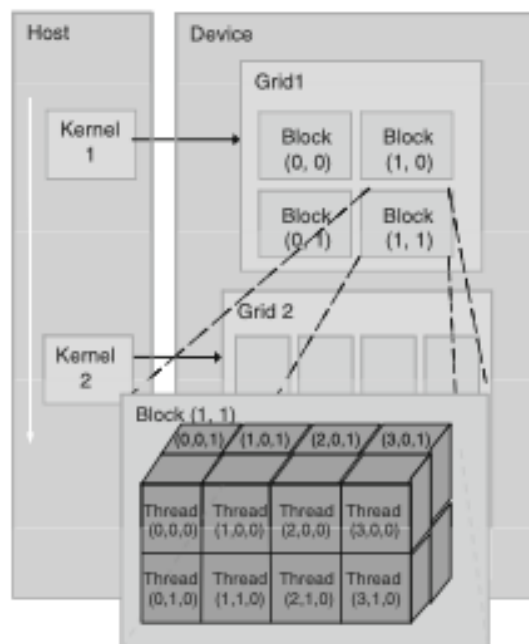
A block is a 3D array of threads (max 512 threads): indexed by threadIdx.x, .y, .z
When launching a kernel, host code chooses grid & block size.
Example:
```
// Setup execution configuration
dim3 dimBlock(Width, Width);
dim3 dimGrid(1, 1);
```

```
// Launch device computation threads!
c MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd, Pd, Width);
```